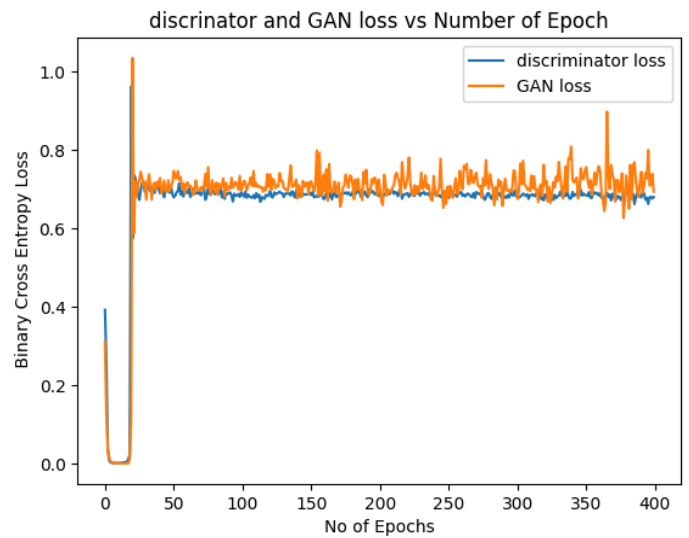





ME 793 Assignment 10

Roll No.: 19D100011

Answer 1:



After doubling the epochs from 200 to 400 and running the notebook again, we observe that the generated images appear to have more refined details and clearer shapes compared to the images generated with the default number of epochs.



ME 793 - Assignment 10

Answer 2

19D100011

```
In [23]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import cv2
```

```
In [24]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from tensorflow.keras import layers
import tensorflow_datasets as tfds
from keras.models import Sequential, Model
from keras.layers import Dense, Flatten, Reshape, Dropout, LeakyReLU, Conv2DTranspose, Conv2D
from keras.optimizers import Adam, SGD
from keras.layers import LeakyReLU, BatchNormalization, Input
```

I used structural similarity index (SSIM) as a quantitative measure of similarity. I used the compare_ssim function from the scikit-image library to compute the SSIM score.

```
In [25]: from skimage.metrics import structural_similarity as ssim
```

```
In [26]: from keras.datasets import mnist
(train_images, y_train), (_, _) = mnist.load_data()
print(len(train_images))
train_images=train_images[(y_train==0) | (y_train==1) | (y_train==5) | (y_train==7)]
train_images=train_images[:5000]

60000
```

```
In [27]: len(train_images)
```


```
Out[27]: 5000
```

```
In [28]: def plot_images(images):
    fig, axes = plt.subplots(5, 5, figsize=(5, 5))

    for i, ax in enumerate(axes.flat):
        ax.imshow(images[i], cmap='gray')
        ax.axis('off')

    plt.show()
```

```
In [29]: plot_images(train_images)
```



```
In [30]: train_images=train_images.astype("float32")/127.5-1 # norm
train_images=train_images.reshape(train_images.shape[0],28,28,1) # now t
```

```
In [31]: def build_generator(latent_dim,input_shape):
generator=Sequential([

    Dense(7*7*128 , input_dim=latent_dim),
    LeakyReLU(alpha=.2),
    Reshape((7, 7, 128)),
    BatchNormalization(),

    Conv2DTranspose(64 , (4,4) , padding="same" , strides=(2,2)),
    LeakyReLU(alpha=.2),
    BatchNormalization(),

    Conv2DTranspose(32 , (4,4) , padding="same" , strides=(2,2)),
    LeakyReLU(alpha=.2),
    BatchNormalization(),

    Conv2DTranspose(1 , (4,4) , padding="same" , activation="tanh"),
])

return generator
```

```
In [32]: def build_discriminator(input_shape):
discriminator=Sequential([

    Conv2D(64 , (3,3) , strides=(2,2) , padding="same" , input_shape=input_shape),
    LeakyReLU(alpha=.2),
    Dropout(rate=.4),

    Conv2D(64 , (3,3) , strides=(2,2) , padding="same" , input_shape=input_shape),
    LeakyReLU(alpha=.2),
    Dropout(rate=.4),

    Flatten(),
```

```

        Dense(1, activation="sigmoid" )

    ])
    return discriminator

```

```

In [33]: def build_gan(generator, discriminator):
          discriminator.trainable = False # Freeze discriminator during GAN training

          gan_input = Input(shape=(latent_dim,))
          x = generator(gan_input)
          gan_output = discriminator(x)

          gan = Model(gan_input, gan_output)
          return gan

```

```

In [34]: latent_dim=100
          image_shape=(28,28,1)

```

```

In [35]: generator = build_generator(latent_dim, image_shape)
          generator.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5))

          discriminator = build_discriminator(image_shape)
          discriminator.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])

          gan = build_gan(generator, discriminator)
          gan.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5))

```

```

In [36]: epochs = 400
          batch_size = 256

```

```

In [37]: def generate_real_samples(dataset, n_samples):
          idx = np.random.randint(0, train_images.shape[0], n_samples)
          X = train_images[idx]
          y = np.ones((n_samples, 1))
          return X, y

          # Generate fake samples using the generator
          def generate_fake_samples(generator, latent_dim, n_samples):
              noise = np.random.normal(0, 1, size=(n_samples, latent_dim))
              X = generator.predict(noise)
              y = np.zeros((n_samples, 1))
              return X, y

          # Generate points in latent space as input for the generator
          def generate_latent_points(latent_dim, n_samples):
              noise = np.random.normal(0, 1, size=(n_samples, latent_dim))
              return noise

```

```

In [38]: discriminator_loss=[]
          gan_loss=[]
          generated_list=[]

```

```

In [39]: def train(g_model, d_model, gan_model, dataset, latent_dim, n_epochs, n_batch):
          bat_per_epo = int(dataset.shape[0] / n_batch)
          half_batch = int(n_batch / 2)
          discriminator_loss=[]
          gan_loss=[]
          generated_list=[]
          # Manually enumerate epochs
          for i in range(n_epochs):
              # Enumerate batches over the training set

```

```

for j in range(bat_per_epo):
    batch_loss_d=[]
    batch_loss_g=[]
    # Get randomly selected 'real' samples
    X_real, y_real = generate_real_samples(dataset, half_batch)

    # Generate 'fake' examples
    X_fake, y_fake = generate_fake_samples(g_model, latent_dim, half_batch)

    # Create training set for the discriminator
    X, y = np.vstack((X_real, X_fake)), np.vstack((y_real, y_fake))

    # Update discriminator model weights
    d_loss, _ = d_model.train_on_batch(X, y)
    batch_loss_d.append(d_loss)
    # Prepare points in latent space as input for the generator
    X_gan = generate_latent_points(latent_dim, n_batch)

    # Create inverted labels for the fake samples
    y_gan = np.ones((n_batch, 1))

    # Update the generator via the discriminator's error
    g_loss = gan_model.train_on_batch(X_gan, y_gan)
    batch_loss_g.append(g_loss)
    gan_loss=np.mean(batch_loss_g)
    gan_loss.append(gan_loss)
    dis_loss=np.mean(batch_loss_d)
    discriminator_loss.append(dis_loss)

    # Summarize loss on this batch

    if (i + 1) % 10 == 0:
        print(f'Epoch {i + 1}/{n_epochs}, Batch {j + 1}/{bat_per_epo}, D Loss {dis_loss}')
    if (i + 1) % 10 == 0:
        num_samples=5
        noise = generate_latent_points(latent_dim, num_samples)
        generated_images = generator.predict(noise)
        generated_list.append(generated_images)

return discriminator_loss, gan_loss, generated_list

```

In [40]: # discriminator_loss, gan_loss, generated_list=train(generator, discriminator, gan,

In [41]: # fig, axes = plt.subplots(len(generated_list), 5, figsize=(5, len(generated_list)))
for i in range(len(generated_list)):
image_list=generated_list[i]
for j in range(5):
axes[i,j].imshow(image_list[j],cmap="gray")
axes[i, j].axis('off')
plt.tight_layout()
plt.show()

In [42]: # plt.plot(discriminator_loss,label="discriminator loss")
plt.plot(gan_loss,label="GAN loss")
plt.legend()
plt.title("discriminator and GAN Loss vs Number of Epoch")
plt.xlabel("No of Epochs")
plt.ylabel(" Binary Cross Entropy Loss")
plt.show()

In [43]: # Select a random input image of digit '7'
input_image = train_images[np.random.randint(0, train_images.shape[0])]

```
input_image = cv2.resize(input_image, (28, 28), interpolation=cv2.INTER_AREA)
input_image = input_image.reshape(1, 28, 28, 1)

# Train the GAN for 200 and 400 epochs
discriminator_loss_200, gan_loss_200, generated_list_200 = train(generator, discrimin
discriminator_loss_400, gan_loss_400, generated_list_400 = train(generator, discrimin
```

Streaming output truncated to the last 5000 lines.

```
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
Epoch 140/400, Batch 19/19, D Loss: 0.688, G Loss: 0.730
1/1 [=====] - 0s 19ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 6ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 6ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 7ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 6ms/step
4/4 [=====] - 0s 4ms/step
```



```

4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 5ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
4/4 [=====] - 0s 3ms/step
Epoch 400/400, Batch 19/19, D Loss: 0.692, G Loss: 0.715
1/1 [=====] - 0s 19ms/step

```

```

In [44]: # Generate images after 200 and 400 epochs
generated_images_200 = generated_list_200[-1]
generated_images_400 = generated_list_400[-1]

# Select a random generated image of digit '7' after 200 and 400 epochs
generated_image_200 = generated_images_200[np.random.randint(0, generated_images_200.shape[0])]
generated_image_400 = generated_images_400[np.random.randint(0, generated_images_400.shape[0])]

# Reshape the images to 28x28 for comparison
input_image = input_image.squeeze()
generated_image_200 = generated_image_200.squeeze()
generated_image_400 = generated_image_400.squeeze()

# Calculate SSIM score
ssim_score_200 = ssim(input_image, generated_image_200, multichannel=True)
ssim_score_400 = ssim(input_image, generated_image_400, multichannel=True)

print(f"SSIM score after 200 epochs: {ssim_score_200:.4f}")
print(f"SSIM score after 400 epochs: {ssim_score_400:.4f}")

# Plot the images
fig, axes = plt.subplots(1, 3, figsize=(15, 5))
axes[0].imshow(input_image, cmap="gray")
axes[0].set_title("Input Image")
axes[0].axis('off')
axes[1].imshow(generated_image_200, cmap="gray")
axes[1].set_title(f"Generated Image (SSIM score: {ssim_score_200:.4f})")
axes[1].axis('off')
axes[2].imshow(generated_image_400, cmap="gray")
axes[2].set_title(f"Generated Image (SSIM score: {ssim_score_400:.4f})")
axes[2].axis('off')
plt.tight_layout()
plt.show()

```

SSIM score after 200 epochs: 0.4563


SSIM score after 400 epochs: 0.4526

<ipython-input-44-e577e0a3c159>:15: FutureWarning: `multichannel` is a deprecated argument name for `structural_similarity`. It will be removed in version 1.0. Please use `channel_axis` instead.

```
ssim_score_200 = ssim(input_image, generated_image_200, multichannel=True)
```

<ipython-input-44-e577e0a3c159>:16: FutureWarning: `multichannel` is a deprecated argument name for `structural_similarity`. It will be removed in version 1.0. Please use `channel_axis` instead.

```
ssim_score_400 = ssim(input_image, generated_image_400, multichannel=True)
```



In []: